

Project 1: Optimizing the Performance of a Pipelined Processor

520021910863, Heda Chen, marcythm@sjtu.edu.cn
520030910287, Xinming Shu, Xinming_Sh@sjtu.edu.cn

April 23, 2022

1 Introduction

This project consists of three parts, which are listed below:

1. In Part A, we need to translate three functions in `sim/misc/examples.c` which are written in C into Y86 assembly program.
2. In Part B, the requirement is to extend the implementation of the given sequential Y86 processor to support the `iaddl` instruction.
3. In Part C, the requirement is to optimize both the given pipelined Y86 processor and `ncopy.y8` to improve the performance of the assembly program. We tried several approaches, such as loop unrolling and rearrange the order of instructions to avoid hazards, and finally achieved a CPE of 9.17.

Additionally, all of the three parts have passed the tests given, and all assembly programs are annotated in detail.

In the process of finishing the project,

1. Chen first wrote all the code, then Shu did the same independently, checked Chen's code for correctness and discussed some details.
2. Shu wrote the main part of the report, Chen revised some details and reformatted the report in \LaTeX .

2 Experiments

2.1 Part A

2.1.1 Analysis

This part is a 'warm-up' session of this project. In this part, we need to translate three functions in `sim/misc/examples.c` which are written in C into Y86 assembly program. The key points and core techniques used are listed as follows:

1. Follow the Y86 calling conventions, e.g. take care to save the value of callee-saved registers (`%ebx`, `%ebp`) in function, pass function arguments using `%edi`, `%esi`, `%edx`...
2. Use stack properly to save caller-saved registers when calling a subroutine.
3. Mimic the functionality of C, divide the code based on functionality with sufficient and clear labels.

For more details, see the comments in the code section.

2.1.2 Code

1. Core part of sum.y

```
1  main:
2      irmovl ele1, %edi    # following the calling conventions, use %rdi
3      call    sum_list     # (here %edi) to store the first argument
4      ret
5  sum_list:
6      irmovl 0, %eax       # initialize `sum`: `sum = 0`
7      jmp     check_condition # jump into loop, and check loop condition first: `while (
                             ls)`
8  L3:
9      mrmovl (%edi), %ecx  # load the value of current element: `tmp = *(ls->val)` (
                             because
10                                     # in Y86 the addition can only be performed between
                                     registers)
11      addl    %ecx, %eax   # add the value to `sum`: `sum += tmp`
12      mrmovl 4(%edi), %edi # move the pointer to the next element: `ls = ls->next`
13  check_condition:
14      andl    %edi, %edi   # check if current pointer is 0: `ls == NULL`
15      jne     L3           # if not, jump to L3
16      ret
```

2. Core part of rsum.y

```
1  main:
2      irmovl ele1, %edi    # following the calling conventions, use %rdi
3      call    rsum_list    # (here %edi) to store the first argument
4      ret
5  rsum_list:
6      andl    %edi, %edi   # check if current pointer is 0 (the end of the list): `if
                             (!ls)`
7      je      L3           # if so, terminate the recursion: `return 0`
8      pushl   %edi         # save the pointer of current element
9      mrmovl 4(%edi), %edi  # move pointer to the next element
10     call    rsum_list     # calculate the sum of the list starting from the next
                             element
11     popl    %edi         # restore the pointer of current element
12     mrmovl (%edi), %ecx   # load the value of current element: `tmp = *(ls->val)` (
                             because
13                                     # in Y86 the addition can only be performed between
                                     registers)
14     addl    %ecx, %eax    # add the value of current element to get the sum of current
                             list
15     ret
16  L3:
17     irmovl 0, %eax       # when terminate the recursion, set counter to 0: `sum = 0`
18     ret
```

3. Core part of copy.y

```
1  main:
2      irmovl src, %edi     # following the calling conventions, use %rdi,
3      irmovl dest, %esi    # %rsi and %rdx (here %edi, %esi and %edx) to
4      irmovl 3, %edx       # store the first three arguments
5      call    copy_block
6      ret
7  copy_block:
8      pushl   %ebx         # save %ebx
9      irmovl 0, %eax       # initialize checksum: `result = 0`
10     jmp     check_condition # jump into loop, and check loop condition first: `while (
                             ls)`
11  L3:
```

```

12    mrmovl (%edi), %ebx # load current value in source block: `val = *src`
13    rmmovl %ebx, (%esi) # store the value into dest block: `*dest = val`
14    xorl   %ebx, %eax    # add current value into checksum: `result ^= val`
15    irmovl 4, %ecx      # use %ecx as a constant 4
16    addl   %ecx, %edi    # move to next value in source block: `src++`
17    addl   %ecx, %esi    # move to next value in dest block: `dest++`
18    irmovl 1, %ecx      # use %ecx as a constant 1
19    subl   %ecx, %edx    # decrease `len`: `len--`
20    check_condition:
21    andl   %edx, %edx    # check if `len` is greater than zero: `while (len > 0)`
22    jg     L3            # if not, continue to loop
23    popl   %ebx         # restore %ebx
24    ret

```

2.1.3 Evaluation

1. For `sum.js` simulating `sum_list()`, as the figure represents, the return value `%eax` is `0xcba`, while none of the callee-saved registers are changed.

```

> ./yas sum.js && ./yis sum.yo
Stopped in 26 steps at PC = 0xb. Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%eax: 0x00000000 0x00000cba
%ecx: 0x00000000 0x00000c00
%esp: 0x00000000 0x00000100

Changes to memory:
0x00f8: 0x00000000 0x0000002f
0x00fc: 0x00000000 0x0000000b

```

Figure 1: the test result of `sum.js`

2. For `rsum.js` simulating `rsum_list()`, as the figure represents, the return value `%eax` is `0xcba`, while none of the callee-saved registers are changed.

```

> ./yas rsum.js && ./yis rsum.yo
Stopped in 37 steps at PC = 0xb. Status 'HLT', CC Z=0 S=0 O=0
Changes to registers:
%eax: 0x00000000 0x00000cba
%ecx: 0x00000000 0x0000000a
%esp: 0x00000000 0x00000100
%edi: 0x00000000 0x0000000c

Changes to memory:
0x00e0: 0x00000000 0x00000044
0x00e4: 0x00000000 0x0000001c
0x00e8: 0x00000000 0x00000044
0x00ec: 0x00000000 0x00000014
0x00f0: 0x00000000 0x00000044
0x00f4: 0x00000000 0x0000000c
0x00f8: 0x00000000 0x0000002f
0x00fc: 0x00000000 0x0000000b

```

Figure 2: the test result of `rsum.js`

3. For `copy.y8` simulating `copy_block()`, as the figure represents, all of the data are correctly copied, and the returned checksum `%eax` is `0xcba`, while none of the callee-saved registers are changed.

```

> ./yas copy.y8 && ./yis copy.y8
Stopped in 45 steps at PC = 0xb, Status 'HLT', CC Z=1 S=0 O=0
Changes to registers:
%eax: 0x00000000 0x00000cba
%ecx: 0x00000000 0x00000001
%esp: 0x00000000 0x00000100
%esi: 0x00000000 0x00000024
%edi: 0x00000000 0x00000018
Changes to memory:
0x0018: 0x00000111 0x0000000a
0x001c: 0x00000222 0x000000b0
0x0020: 0x00000333 0x00000c00
0x00f8: 0x00000000 0x0000003b
0x00fc: 0x00000000 0x0000000b

```

Figure 3: the test result of `copy.y8`

2.2 Part B

2.2.1 Analysis

In this part, we should extend the sequential Y86 processor by modifying `seq-full.hcl` to support `iaddl` instruction. The key points are listed as follows:

1. Understand the processor's logic and take a good command of the syntax of HCL.
2. Determine which signals should be modified to implement `iaddl` instruction.

The process of determining which signals should be modified is as follows:

1. In Fetch Stage, `IIADDL` should be added into `instr_valid`, `need_regid` and `need_valC` because it is a valid instruction, and also need `regid` and `valC`.
2. In Decode Stage, when `icode` is `IIADDL`, `srcB` is from `rB` since the second operand of `iaddl` is a register, `dstE` (where the result from ALU is passed towards) is `rB` since `iaddl imm, rB` means `rB += imm` (`rB` is updated).
3. In Execute Stage, add `IIADDL` into the choices region of `set_cc` since `iaddl` operation involves ALU operation which will set conditional codes.
4. In Execute Stage, when `icode` is `IIADDL`, `aluA` (the first op) is `valC` (the immediate in the instruction) since `iaddl imm, rB` means the first op is `imm` (`valC`).
5. In Execute Stage, when `icode` is `IIADDL`, `aluB` (the second op) is `valB` (the value of the second register that is read) for the same reason above.

2.2.2 Code

Here only the modified parts of the code are listed.

```

1 bool instr_valid = icode in
2   { INOP, IHALT, IRRMOVL, IIRMOVL, IRMMOVL, IMRMOVL,
3     IOPL, IJXX, ICALL, IRET, IPUSHL, IPOPL, IIADDL };
4

```

```

5 # Does fetched instruction require a regid byte?
6 bool need_regids =
7     icode in { IRRMOVL, IOPL, IPUSHL, IPOPL,
8               IIRMOVL, IRMMOVL, IMRMOVL, IIADDL };
9
10 # Does fetched instruction require a constant word?
11 bool need_valC =
12     icode in { IIRMOVL, IRMMOVL, IMRMOVL, IJXX, ICALL, IIADDL };
13
14 ## What register should be used as the B source?
15 int srcB = [
16     icode in { IOPL, IRMMOVL, IMRMOVL, IIADDL } : rB;
17     icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;
18     1 : RNONE; # Don't need register
19 ];
20
21 ## What register should be used as the E destination?
22 int dstE = [
23     icode in { IRRMOVL } && Cnd : rB;
24     icode in { IIRMOVL, IOPL, IIADDL } : rB;
25     icode in { IPUSHL, IPOPL, ICALL, IRET } : RESP;
26     1 : RNONE; # Don't write any register
27 ];
28
29 ## Select input A to ALU
30 int aluA = [
31     icode in { IRRMOVL, IOPL } : valA;
32     icode in { IIRMOVL, IRMMOVL, IMRMOVL, IIADDL } : valC;
33     icode in { ICALL, IPUSHL } : -4;
34     icode in { IRET, IPOPL } : 4;
35     # Other instructions don't need ALU
36 ];
37
38 ## Select input B to ALU
39 int aluB = [
40     icode in { IRMMOVL, IMRMOVL, IOPL, ICALL,
41               IPUSHL, IRET, IPOPL, IIADDL } : valB;
42     icode in { IRRMOVL, IIRMOVL } : 0;
43     # Other instructions don't need ALU
44 ];
45
46 ## Should the condition codes be updated?
47 bool set_cc = icode in { IOPL, IIADDL };

```

2.2.3 Evaluation

1. First, we test out solution on the Y86 program `asumi.js`. Here the result is omitted because it takes too much space and is not good for the layout of the report.
2. Then we retest the solution using the benchmark programs. Here the result is omitted because of the same reason as above.
3. Finally we test the implementation of `iaddl` with regression tests. As the figure shows, our implementation passes all the ISA checks.

2.3 Part C

2.3.1 Analysis

In this part, we were asked to speed up the program `ncopy.js` as much as possible by optimizing both the program itself and the pipelined Y86 processor's implementation. The key points and core techniques are listed as follows:

```

> cd ../ptest; make SIM=../seq/ssim TFLAGS=-i
./optest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 58 ISA Checks Succeed
./jtest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 96 ISA Checks Succeed
./ctest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 22 ISA Checks Succeed
./htest.pl -s ../seq/ssim -i
Simulating with ../seq/ssim
All 756 ISA Checks Succeed

```

Figure 4: Part B Regression test

1. Modify `pipe-full.hcl` to support `iaddl` instruction. This part is similar to Part B, so I will not go into details.
2. Use `iaddl` instruction to replace the load-add case, e.g. replace `irmovl 1, %edi; addl %edi, %eax` with `iaddl 1, %eax`.
3. Perform loop unrolling to reduce the overhead of frequently modifying and comparing the subscript in the loop. Here we choose 7-way loop unrolling, as this maximizes the exploitation of the binary search technique used later.
4. Avoid load-and-use hazards. When a load-and-use hazard occurs, it will waste a clock cycle, which can be avoided and utilized by reordering instructions. Furthermore, we use two registers (`%esi`, `%edi`) to store the variable `val` in turn, loading them separately ahead of time.
5. Use binary search for remaining elements. For large inputs, it's better to unroll the loops for more ways. However, for small inputs, it is important to choose a good way to handle the remaining elements. For parts smaller than the number of ways, the easiest way is to write another loop for them. But there is another approach that performs better, that is, jumping to different positions for different number of remaining elements. Since Y86 does not support relative jump instructions, we use binary search to get the correct jump destination for each case.
Additionally, here we choose to implement a binary search of depth 2, thus the maximum distinguishable number is $2^2 + 2^1 + 2^0 = 7$, so we use 7-way loop unrolling as above.

For more details, see the comments in the code section.

2.3.2 Code

1. 7-way loop unrolling part

```

1 #####
2 # You can modify this portion
3 xorl %eax,%eax # count = 0;
4 iaddl -6, %edx # len <= 6?
5 jle S06       # if so, goto S06
6
7 L0:
8 mrmovl (%ebx), %esi # tmp0 = src[0]
9 mrmovl 4(%ebx), %edi # tmp1 = src[1]
10 andl %esi, %esi     # if tmp0 <= 0
11 rmmovl %esi, (%ecx) # dst[0] = tmp0
12 jle L1             # skip increasing counter
13 iaddl 1, %eax       # else counter++

```

```

14 L1:
15     mrmovl 8(%ebx), %esi # tmp0 = src[2]
16     andl %edi, %edi      # if tmp1 <= 0
17     rmmovl %edi, 4(%ecx) # dst[1] = tmp1
18     jle L2               # skip increasing counter
19     iaddl 1, %eax        # else counter++
20 L2:
21     mrmovl 12(%ebx), %edi # tmp1 = src[3]
22     andl %esi, %esi      # if tmp0 <= 0
23     rmmovl %esi, 8(%ecx) # dst[2] = tmp0
24     jle L3               # skip increasing counter
25     iaddl 1, %eax        # else counter++
26 L3:
27     mrmovl 16(%ebx), %esi # tmp0 = src[4]
28     andl %edi, %edi      # if tmp1 <= 0
29     rmmovl %edi, 12(%ecx) # dst[3] = tmp1
30     jle L4               # skip increasing counter
31     iaddl 1, %eax        # else counter++
32 L4:
33     mrmovl 20(%ebx), %edi # tmp1 = src[5]
34     andl %esi, %esi      # if tmp0 <= 0
35     rmmovl %esi, 16(%ecx) # dst[4] = tmp0
36     jle L5               # skip increasing counter
37     iaddl 1, %eax        # else counter++
38 L5:
39     mrmovl 24(%ebx), %esi # tmp0 = src[6]
40     andl %edi, %edi      # if tmp1 <= 0
41     rmmovl %edi, 20(%ecx) # dst[5] = tmp1
42     jle L6               # skip increasing counter
43     iaddl 1, %eax        # else counter++
44 L6:
45     andl %esi, %esi      # if tmp0 <= 0
46     rmmovl %esi, 24(%ecx) # dst[6] = tmp0
47     jle ChkCond          # skip increasing counter
48     iaddl 1, %eax        # else counter++
49 ChkCond:
50     iaddl 28, %ebx        # src += 7 * sizeof(Byte)
51     iaddl 28, %ecx        # dst += 7 * sizeof(Byte)
52     iaddl -7, %edx        # len -= 7
53     jg L0                # if len > 0, goto L0

```

2. Binary search for finding the number of remaining elements

```

1  S06:                # len in (x-6: [0, 6] -> [-6, 0])
2  iaddl 3, %edx        # (x-6: [0, 6] -> [-6, 0]) => (x-3: [0, 6] -> [-3, 3])
3  jg S46               # len-3 > 0, len in [4, 6]
4  je R3                # len-3 == 0, len = 3
5  S02:                # len in (x-3: [0, 2] -> [-3, -1])
6  iaddl 2, %edx        # (x-3: [0, 2] -> [-3, -1]) => (x-1: [0, 2] -> [-1, 1])
7  jl R0                # len-1 < 0, len = 0
8  je R1                # len-1 == 0, len = 1
9  jmp R2               # len-1 > 0, len = 2
10 S46:                # len in (x-3: [4, 6] -> [1, 3])
11 iaddl -2, %edx        # (x-3: [4, 6] -> [1, 3]) => (x-5: [4, 6] -> [-1, 1])
12 jl R4                # len-5 < 0, len = 4
13 je R5                # len-5 == 0, len = 5

```

3. Unrolling of remaining loops

```

1  R6:
2  mrmovl 20(%ebx), %esi # tmp = src[6]
3  andl %esi, %esi      # if tmp <= 0
4  rmmovl %esi, 20(%ecx) # dst[6] = tmp

```

```

5   jle R5          # skip increasing counter
6   iaddl 1, %eax   # else counter++
7   R5:
8   mrmovl 16(%ebx), %esi # tmp = src[5]
9   andl %esi, %esi   # if tmp <= 0
10  rmmovl %esi, 16(%ecx) # dst[5] = tmp
11  jle R4          # skip increasing counter
12  iaddl 1, %eax   # else counter++
13  R4:
14  mrmovl 12(%ebx), %esi # tmp = src[4]
15  andl %esi, %esi   # if tmp <= 0
16  rmmovl %esi, 12(%ecx) # dst[4] = tmp
17  jle R3          # skip increasing counter
18  iaddl 1, %eax   # else counter++
19  R3:
20  mrmovl 8(%ebx), %esi # tmp = src[3]
21  andl %esi, %esi   # if tmp <= 0
22  rmmovl %esi, 8(%ecx) # dst[3] = tmp
23  jle R2          # skip increasing counter
24  iaddl 1, %eax   # else counter++
25  R2:
26  mrmovl 4(%ebx), %esi # tmp = src[2]
27  andl %esi, %esi   # if tmp <= 0
28  rmmovl %esi, 4(%ecx) # dst[2] = tmp
29  jle R1          # skip increasing counter
30  iaddl 1, %eax   # else counter++
31  R1:
32  mrmovl (%ebx), %esi # tmp = src[1]
33  andl %esi, %esi   # if tmp <= 0
34  rmmovl %esi, (%ecx) # dst[1] = tmp
35  jle R0          # skip increasing counter
36  iaddl 1, %eax   # else counter++
37  R0:

```

2.3.3 Evaluation

1. Run the given correctness.pl script to check correctness:

```

1  > ./correctness.pl
2  Simulating with instruction set simulator yis
3  ncopy
4  0      OK
5  1      OK
6  2      OK
7  3      OK
8  # 61 lines omitted
9  64     OK
10 128     OK
11 192     OK
12 256     OK
13 68/68 pass correctness test

```

2. Run benchmark.pl script to check performance, and get a CPE of 9.17:

```

1  > ./benchmark.pl
2  ncopy
3  0      39
4  1      47      47.00
5  2      58      29.00
6  3      53      17.67
7  # 59 lines omitted
8  62     422      6.81
9  63     424      6.73
10 64     435      6.80

```



```
11 Average CPE 9.17
12 Score 60.0/60.0
```

3 Conclusion

3.1 Problems

Several obstacles we encountered are as follows:

1. Most of the information available on the Internet is about the newer 64-bit Y86, which has some minor differences from the 32-bit Y86 used in this experiment, causing some troubles. For example, when Chen was writing some sample programs to help him understand Y86 ISA, the assembler kept reporting an confusing error `Invalid line`. After checking the program several times, he realized that he should change `xxxq` to `xxxl`, as the older CS:APP books used.
2. The assembler gives very little information about syntax error, which is hardly helpful for debugging. Chen often spends quite a long time, looking for where exactly the syntax error is. The example is the same as above.
3. When Chen started trying Part C, he first attempted to replace the load-add case with the `iaddl` instruction. Surprisingly, the benchmark script gives a CPE of 2 at this point (of course, the correctness script gave the positive result). Chen checked for a long time and even suspected that there was a problem with the benchmark script (and actually rewrote it). Finally he found out that the `iaddl` instruction was not implemented in `pipe-full.hcl`. After migrating the implementation from Part B, the CPE was back to normal level. (In fact this question is still unsolved: why could that version pass the correctness test? Is there any bug?)

3.2 Achievements

1. First of all, the most remarkable advantage of our solution should be the readability of the code. In every assembly program, almost every instruction is followed by a comment explaining what it does. (For HCL files, the original comments were clear enough, so we didn't add comment to them.)
2. Also, the good code readability is beneficial for collaboration with partner. We use GitHub for code synchronization, and discussed details thoroughly when writing the report.
3. With respect to the performance of our solution, we have optimized the Part C code to maximize the use of the above techniques under limited conditions (depth = 2), achieving the CPE of 9.17. In fact Chen also investigated how to implement branch prediction, but did not attempt it after discovering that loop unrolling was sufficient enough to obtain a low CPE because of the hassle and lack of time.
4. In this experiment we successfully applied what we have learned in class and gained a deeper understanding.